

6. Monitoring

Introduction to PostgreSQL



PostgreSQL

AGENDA

- Configuration Monitoring
- Base Performance Monitoring
- Advanced Performance Monitoring
- Tools



MAIN CONFIG FILES

- There are three main configuration files that present the starting point for any configuration:
 - postgresql.conf is the main cluster configuration file
 - *contains all the data required to start the cluster, set up processes (as WAL senders) and logging, and configure how the cluster will accept connections*
 - postgresql.auto.conf automatically generated and edited by the cluster itself
 - *contains parameters changed by the superuser from within the cluster.*
 - pg_hba.conf used to allow or deny the client connections to the cluster.
- A parameter can be defined multiple times but only the last definition found is used
 - Every parameter has a default value used if it is not defined in the config specs

RUNNING CONFIG

- You can inspect the running configuration by running a query on *pg_settings*
 - You will do this in more detail in the lab.

```
rod=# SELECT name, setting || ' ' || unit as current_value,
min_val, max_val, boot_val, reset_val
from pg_settings;
```

name	current_value	min_val	max_val	boot_val	reset_val
allow_in_place_tablespaces				off	off
allow_system_table_mods				off	off
application_name					psql
archive_cleanup_command					
archive_command					
archive_library					
archive_mode				off	off
archive_timeout	0 s	0	1073741823	0	0
array_nulls				on	on
authentication_timeout	60 s	1	600	60	60
autovacuum				on	on
autovacuum_analyze_scale_factor		0	100	0.1	0.1
autovacuum_analyze_threshold		0	2147483647	50	50
autovacuum_freeze_max_age		100000	2000000000	200000000	2000000000
autovacuum_max_workers		1	262143	3	3
autovacuum_multixact_freeze_max_age		10000	2000000000	400000000	4000000000
autovacuum_naptime	60 s	1	2147483	60	60
autovacuum_vacuum_cost_delay	2 ms	-1	100	2	2
autovacuum_vacuum_cost_limit		-1	10000	-1	-1
autovacuum_vacuum_insert_scale_factor		0	100	0.2	0.2
autovacuum_vacuum_insert_threshold		-1	2147483647	1000	1000
autovacuum_vacuum_scale_factor		0	100	0.2	0.2
autovacuum_vacuum_threshold		0	2147483647	50	50
autovacuum_work_mem	-1 kB	-1	2147483647	-1	-1
backend_flush_after	0 8kB	0	256	0	0
backslash_quote				safe_encoding	safe_encoding
backtrace_functions					
bgwriter_delay	200 ms	10	10000	200	200
bgwriter_flush_after	64 8kB	0	256	64	64
bgwriter_lru_maxpages		0	1073741823	100	100
bgwriter_lru_multiplier		0	10	2	2

RUNNING CONFIG

- The pg_settings contains for each parameter
 - The file and the line number from where a parameter has been loaded.
 - Used to find where a configuration parameter has been set
 - In the screenshot, only the parameters set in a file are listed.

```
rod=# SELECT name, setting as current_value,
sourcefile, sourceline, pending_restart
from pg_settings
where sourcefile is Not null;
```

name	current_value	sourcefile	sourceline	pending_restart
DateStyle	ISO, YMD	/var/lib/pgsql/data/postgresql.conf	715	f
default_text_search_config	pg_catalog.english	/var/lib/pgsql/data/postgresql.conf	741	f
dynamic_shared_memory_type	posix	/var/lib/pgsql/data/postgresql.conf	153	f
lc_messages	en_CA.UTF-8	/var/lib/pgsql/data/postgresql.conf	731	f
lc_monetary	en_CA.UTF-8	/var/lib/pgsql/data/postgresql.conf	733	f
lc_numeric	en_CA.UTF-8	/var/lib/pgsql/data/postgresql.conf	734	f
lc_time	en_CA.UTF-8	/var/lib/pgsql/data/postgresql.conf	735	f
listen_addresses	*	/var/lib/pgsql/data/postgresql.conf	60	f
log_filename	postgresql-%a.log	/var/lib/pgsql/data/postgresql.conf	465	f
log_rotation_age	1440	/var/lib/pgsql/data/postgresql.conf	469	f
log_timezone	America/Winnipeg	/var/lib/pgsql/data/postgresql.conf	603	f
log_truncate_on_rotation	on	/var/lib/pgsql/data/postgresql.conf	474	f
logging_collector	on	/var/lib/pgsql/data/postgresql.conf	457	f
max_connections	100	/var/lib/pgsql/data/postgresql.conf	65	f
max_wal_size	1024	/var/lib/pgsql/data/postgresql.conf	247	f
min_wal_size	80	/var/lib/pgsql/data/postgresql.conf	248	f
port	5432	/var/lib/pgsql/data/postgresql.conf	64	f
shared_buffers	16384	/var/lib/pgsql/data/postgresql.conf	130	f
TimeZone	America/Winnipeg	/var/lib/pgsql/data/postgresql.conf	717	f

(19 rows)

CONFIG SETTINGS

- For each of the parameters we looked at in the the previous slides, we can use the SHOW command.
 - We can also use the SET command to override for the current session
 - Or ALTER SYSTEM SET to permanently override the parameter

```
rod=# SHOW work_mem ;
work_mem
-----
4MB
(1 row)

rod=# SET work_mem = '8MB';
SET
rod=# SHOW work_mem ;
work_mem
-----
8MB
(1 row)

rod=# ALTER SYSTEM SET work_mem = '8MB';
ALTER SYSTEM
rod=# SHOW work_mem ;
work_mem
-----
8MB
(1 row)
```

LOCATING CONFIG

- Because the parameters can be set in different files, the table `pg_file_settings` can show exactly where the current parameter value was applied from.
 - Because we overrode the value of `work_mem` with the `ALTER SYSTEM` command, we see that here

```
rod=# SELECT name, setting, sourcefile, sourceline,
applied
from pg_file_settings
where name = 'work_mem';
   name   | setting |          sourcefile          | sourceline | applied
-----+-----+-----+-----+-----
work_mem | 8MB     | /var/lib/pgsql/data/postgresql.auto.conf |          3 | t
(1 row)
```


LOCATING CONFIG

- We can also find out which parameters were not successfully applied but checking the applied flag and any possible errors
 - In this case, no parameters were not applied
 - The error field may contain a message describing the error that prevented the parameter from not being set, like an invalid or out of range value;


```
rod=# SELECT *
FROM pg_file_settings
WHERE applied = false;
 sourcefile | sourceline | seqno | name | setting | applied | error
-----+-----+-----+-----+-----+-----+-----
(0 rows)
```


CONFIGURATION CONTEXTS

- Each configuration parameter belongs to a context
 - This is a group that defines when a change to the parameter can be applied.
 - Some parameters can be changed and take effect during the cluster's life cycle.
 - Others cannot and require the cluster to be restarted
- Contexts are:
 - internal: Value depends on the PostgreSQL source code and is established at compile time, and can only be changed by recompiling.
 - postmaster: The cluster must be restarted a change to be applied.
 - sighup: Changes when given a hang-up signal, usually a reload of the operating system service.
 - superuser-backend and backhand: Changes will be applied to both the client and administrator from the next connection of either type.
 - user and superuser: These changes will be applied immediately to the current connection,

CONFIGURATION GENERATORS

- Configuration generators are used to build a base configuration for a cluster that can be tuned.
 - We will not be exploring these in detail.
 - The lab explores a commonly used one 'pgconfig.org'

**PostgreSQL**
Configuration Builder

HomeContributeDocumentation

Server

Operating system
GNU/Linux Based

Architecture
64 Bits (x86-64)

Storage type
SSD Storage

Number of CPUs
2

Total Memory (GB)
4

Max connections
100

Database

Application profile
General web applications

PostgreSQL Version
16 (Latest)

Profile ComparisonExport Config

Memory Configuration

	Default Value	WEB	OLTP	DW	MIXED	DESKTOP
> shared_buffers	128MB	1GB	1GB	1GB	512MB	256MB
> effective_cache_size	4GB	3GB	3GB	3GB	2GB	614MB
> work_mem	4MB	10MB	14MB	20MB	4MB	839kB
> maintenance_work_mem	64MB	205MB	205MB	205MB	102MB	41MB

Checkpoint Related Configuration

LAB 6-1

- The lab description and documentation is in the Lab directory in the class repository



BASIC OS MONITORING

- Basic monitoring in PostgreSQL focuses on ensuring the general health and performance of the database.
- The key areas include checking system resources, database connectivity, instance availability, and basic performance metrics.
- Instance Availability and Connection Monitoring
 - `pg_isready`, simple command-line tool used to check the availability of a PostgreSQL instance. Equivalent to a "ping" for the database.
- System Resource Monitoring
 - CPU and Memory: Monitor CPU usage, memory usage, and swap space to ensure the database has sufficient resources.
 - Disk I/O: Monitor read/write latency, IOPS, and throughput using tools like `iostat`, `vmstat`, or system-specific monitoring solutions.
 - Disk Space: Regularly check disk space to prevent out-of-space errors, which can severely impact PostgreSQL, as it relies heavily on disk for data storage and temporary operations.

SYSTEM TOOLS MONITORING

- Generally PostgreSQL instances run on Unix
 - There are a number of Unix system monitoring tools that can be used
 - These treat PostgreSQL like any other process
 - Useful for monitoring resource usage
- 'ps' - monitors system processes, checks the status of PostgreSQL processes
 - `"ps aux | grep postgres"`
- 'top and htop' - monitor real-time system performance, including CPU, memory usage being for each running processes.
- 'vmstat' - Monitor system performance metrics, including CPU, memory, and I/O statistics.
 - Check for I/O bottlenecks and memory swapping that could impact PostgreSQL performance.

SYSTEM TOOLS MONITORING

- 'iostat' - monitor disk I/O statistics
 - Identify disk read/write patterns and potential I/O bottlenecks that can affect database performance.
- 'netstat' - Monitor network connections and statistics
- 'pg_top' - A PostgreSQL-specific tool based on top for monitoring active sessions, queries, and performance statistics.
 - View running queries, resource consumption by each session, and overall database activity.

SYSTEM MONITORING

- Graphical OS monitoring tools like Grafana and Prometheus are generally employed
 - Allow for various sort of alerts
 - Not covered in this course



BASIC PERFORMANCE MONITORING

- Database Logs

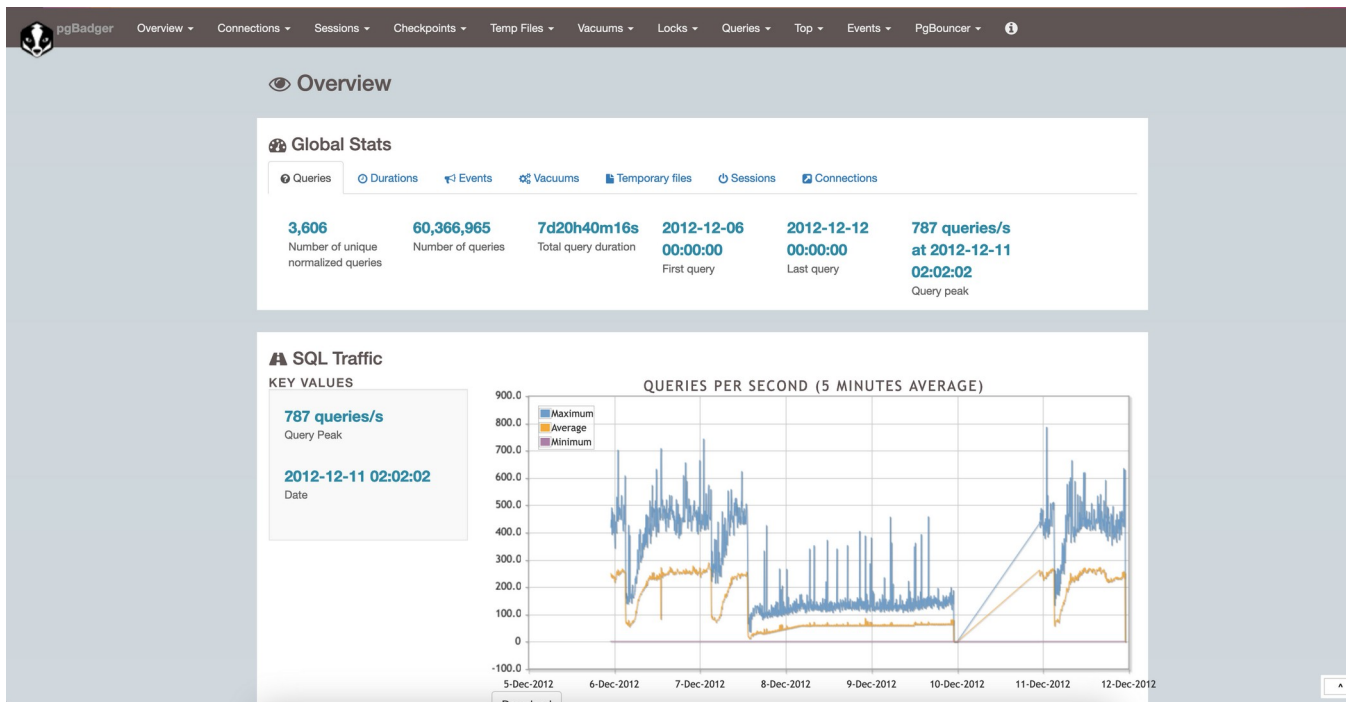
- Log Files: PostgreSQL logs important events such as connection attempts, errors, warnings, and slow queries.
- Configuration: Set appropriate logging levels in postgresql.conf (log_min_duration_statement, log_error_verbosity, etc.) to capture relevant events without overwhelming the system with logs.
- Log Analysis: Use tools like pgBadger or custom scripts to parse and analyze logs for insights into errors, performance issues, and suspicious activities.

- Basic Performance Metrics

- Tools: pg_stat_activity, pg_stat_database, pg_stat_user_tables.
- Key Metrics: Active sessions and connections.
- Database-wide statistics such as transaction counts, number of deadlocks, and cache hit ratios.
- Table-level statistics like sequential scans, index scans, and tuples read/returned.
- Usage: These views provide an overall view of database activity and can be queried regularly or integrated into a monitoring dashboard.

PGBADGER

- Log analysis tool
 - Officially supported by PostgreSQL
 - Provides graphic and quantitative reports on activities.
 - Link is in the notes



COLLECTING STATS

- The cluster collects information about activities by means of the statistic collector
 - A dedicated process that collects cluster-wide information
 - Statistics are not in real time, they are updated at least every 500 milliseconds by backend processes
 - Statistics within a transaction block are “frozen” and cannot be collected until the transaction terminates
 - Statistics are kept across shutdowns and restarts, but are lost after a crash and are reset.
 - Individual database stats can be reset with `pg_stat_reset()`

QUERIES AND SESSIONS

- *pg_stat_activity*
catalog records every
back end process
active in the cluster
 - Only reports the last
executed query from a
session or connection.
 - Not updated until a
new statement is
executed.

```
state
backend_xid
rod=# \d pg_stat_activity
```

View "pg_catalog.pg_stat_activity"				
Column	Type	Collation	Nullable	Default
datid	oid			
datname	name			
pid	integer			
leader_pid	integer			
usesysid	oid			
username	name			
application_name	text			
client_addr	inet			
client_hostname	text			
client_port	integer			
backend_start	timestamp with time zone			
xact_start	timestamp with time zone			
query_start	timestamp with time zone			
state_change	timestamp with time zone			
wait_event_type	text			
wait_event	text			
state	text			
backend_xid	xid			
backend_xmin	xid			
query_id	bigint			
query	text			
backend_type	text			

QUERIES AND SESSIONS

```
rod=# SELECT username, datname, client_addr, application_name,  
backend_start, query_start,  
state, backend_xid, query  
FROM pg_stat_activity;
```

```
query  
-[ RECORD 2 ]-----+-----  
username      | postgres  
datname       |  
client_addr   |  
application_name |  
backend_start  | 2024-11-20 08:16:17.546575-06  
query_start   |  
state         |  
backend_xid    |  
query         |  
-[ RECORD 3 ]-----+-----  
username      | rod  
datname       | rod  
client_addr   |  
application_name | psql  
backend_start  | 2024-12-02 19:31:07.463686-06  
query_start   | 2024-12-02 19:37:45.067357-06  
state         | active  
backend_xid    |  
query         | SELECT username, datname, client_addr, application_name,+  
              | backend_start, query_start,                               +  
              | state, backend_xid, query                                +  
              | FROM pg_stat_activity;  
-[ RECORD 4 ]-----+-----  
username      |
```

LOCKING

- MVCC ensures that only one transaction can modify a particular resource at a time (e.g., a row or table)
 - MVCC allows readers to access older versions of data
 - Does not inherently prevent multiple writers from conflicting.
 - Locks (such as row-level locks like `SELECT ... FOR UPDATE`) prevent such write conflicts.
 - Locks are required to control schema modifications, such as altering a table, creating an index, or dropping a table.
 - Locks manage access to non-database resources or custom application logic using advisory locks, which MVCC does not support.
 - Locks help enforce certain constraints, such as foreign key constraints or unique constraints, by preventing conflicting operations.
 - And other use cases.

MOINTORING LOCKS

- The pg_locks special catalog records any locks that are acquired by different transactions and statements.
 - This catalog allows the system administrator to identify possible bottlenecks
 - Often queried by joining with pg_stat_activity to get more detailed information

```
backend_start | 2024-11-26 09:15:23.091624-05
student=# SELECT a.username, a.application_name, a.datname, a.query,
student=# l.granted, l.mode
student=# FROM pg_locks l
student=# JOIN pg_stat_activity a ON a.pid = l.pid;
-[ RECORD 1 ]-----+-----
username          | student
application_name   | psql
datname            | student
query              | SELECT a.username, a.application_name, a.datname, a.query, +
                  | l.granted, l.mode                                         +
                  | FROM pg_locks l                                         +
                  | JOIN pg_stat_activity a ON a.pid = l.pid;
granted            | t
mode               | AccessShareLock
-[ RECORD 2 ]-----+-----
```


MONITORING DATABASES

- Detailed information about databases are in the pg_stat_database catalog.
 - Provides information about COMMIT and ROLLBACK transactions, deadlocks, and conflicts.
 - If you see the numbers of COMMITs and ROLLBACKs grow quickly, this may be due to an application error or clients making errors in a database forcing rollbacks

```
student=# \x
Expanded display is off.
student=# SELECT datname, xact_commit, xact_rollback, blks_read,
student=# conflicts, deadlocks,
student=# tup_fetched, tup_inserted, tup_updated, tup_deleted, stats_reset
student=# FROM pg_stat_database;
 datname | xact_commit | xact_rollback | blks_read | conflicts | deadlocks | tup_fetched | tup_inserted | tup_updated | tup_deleted | stats_res
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
---
postgres |         0 |          0 |       158 |          0 |          0 |       91407 |         38 |          8 |          5 |
template1 |      150825 |           2 |      1962 |          0 |          0 |      178499 |        821 |        111 |         91 |
template0 |          0 |           0 |          0 |          0 |          0 |          0 |          0 |          0 |          0 |
student   |       6895 |          19 |          0 |          0 |          0 |       55138 |         67 |          1 |         31 |
zorgo     |       1980 |           3 |          0 |          0 |          0 |       21222 |         19 |          0 |          0 |
(6 rows)
```

ADVANCED MONITORING

- Advanced Performance Monitoring
 - `pg_stat_bgwriter`: Provides insights into background writer activity, checkpoints, and buffer management, which are critical for understanding I/O performance.
 - Query Optimization: Utilize `EXPLAIN` and `EXPLAIN ANALYZE` to review query plans and execution times, identifying areas where indexes, rewrites, or configuration changes could improve performance.
- Resource Contention and Wait Event Analysis
 - Wait Events: PostgreSQL has a wait event monitoring system that tracks what processes are waiting on, similar to Oracle's wait events.
 - Blocking and Deadlock Detection: Identify and resolve lock contention and deadlocks using `pg_locks` and `pg_blocking_pids()` functions.
 - Buffer and Cache Analysis: Monitor buffer usage and cache hit ratios through `pg_buffercache` and `pg_stat_database`. Low cache hit ratios can indicate inefficient queries or inadequate memory allocation.

ADVANCED MONITORING

- Replication and High Availability Monitoring
 - Streaming Replication: Monitor replication status using views like `pg_stat_replication`, which shows the state of replication, lag, and connection status of replicas.
 - Failover and Recovery: Tools like `repmgr` or `Patroni` can be used to manage and monitor high availability setups, providing alerts and automation for failover scenarios.
- System-Level Metrics Integration
 - Prometheus and Grafana: Use Prometheus with exporters like `node_exporter` and `postgres_exporter` for detailed time-series data on system and database performance.
 - Detailed Metrics: Monitor CPU, memory, disk I/O, and network metrics alongside database-specific metrics for a holistic view of system health.
- Alerting and Automated Responses
 - Threshold-Based Alerts: Set up alerts for critical metrics such as high CPU usage, long-running queries, replication lag, and disk space.
 - Automated Responses: Implement automated actions in response to certain alerts, such as killing runaway queries, initiating failovers, or adjusting configuration parameters.

ADVANCED MONITORING

- Replication and High Availability Monitoring
 - Streaming Replication: Monitor replication status using views like `pg_stat_replication`, which shows the state of replication, lag, and connection status of replicas.
 - Failover and Recovery: Tools like `repmgr` or `Patroni` can be used to manage and monitor high availability setups, providing alerts and automation for failover scenarios.
- System-Level Metrics Integration
 - Prometheus and Grafana: Use Prometheus with exporters like `node_exporter` and `postgres_exporter` for detailed time-series data on system and database performance.
 - Detailed Metrics: Monitor CPU, memory, disk I/O, and network metrics alongside database-specific metrics for a holistic view of system health.
- Alerting and Automated Responses
 - Threshold-Based Alerts: Set up alerts for critical metrics such as high CPU usage, long-running queries, replication lag, and disk space.
 - Automated Responses: Implement automated actions in response to certain alerts, such as killing runaway queries, initiating failovers, or adjusting configuration parameters.

ADVANCED MONITORING

- Advanced Log Analysis
 - Log Aggregation and Analysis: Use tools like ELK Stack (Elasticsearch, Logstash, Kibana) or Splunk to aggregate and analyze PostgreSQL logs for patterns, anomalies, and trends.
 - Custom Alerts from Logs: Set up alerts based on specific log entries, such as repeated connection failures, frequent checkpoint warnings, or slow queries exceeding a defined threshold.
- Key Differences from Oracle Monitoring
 - Schema and Object Ownership: PostgreSQL's schema design and object ownership can influence monitoring strategies, as it has a more granular permissions model compared to Oracle.
 - Autovacuum and Vacuum Monitoring: Unlike Oracle's automated segment management, PostgreSQL relies on autovacuum processes to reclaim space, which needs to be closely monitored and tuned to avoid performance degradation.
 - Configuration Flexibility: PostgreSQL allows extensive tuning via configuration files (postgresql.conf), requiring regular review and adjustment based on monitored performance data.

MONITORING TOOLS

- Sampling of the most common tools
 - Links to the tools are in the Notes file in the repository
- *PdAdmin*: A popular open-source tool that provides a graphical interface for managing PostgreSQL databases. It includes built-in monitoring and graphing tools.
- *pg_stat_statements*: A PostgreSQL extension that tracks execution statistics of all SQL statements executed by a server.
 - Available within PostgreSQL; it needs to be enabled by adding it to the `shared_preload_libraries` in the `postgresql.conf` file.
- *PgBadger*: A fast PostgreSQL log analyzer that generates detailed reports on performance based on log files.

MONITORING TOOLS

- *pg_top*: Similar to the Unix `top` command but for PostgreSQL, providing a real-time view of database processes, queries, and statistics.
- *PostgreSQL Exporter for Prometheus*: Collects PostgreSQL metrics and exports them to Prometheus for monitoring and alerting.
- *Pgmetrics*: A command-line tool that collects various statistics and configurations from a running PostgreSQL server and displays them in a detailed report format.
- *Percona*: A comprehensive monitoring tool that supports PostgreSQL and offers insights into database performance with advanced dashboards and alerts.
- *pg_stat_kcache*: A PostgreSQL extension that provides statistics on CPU and I/O usage for all SQL statements.

MONITORING TOOLS

- *TimescaleDB*: A time-series database based on PostgreSQL that includes additional features for monitoring and managing time-series data.
- *pgCloo*: A PostgreSQL clusters performance monitoring tool that collects, stores, and visualizes statistics from PostgreSQL and the operating system.

LAB 6-2

- The lab materials and instructions are in the repository



End Module



PostgreSQL