

7. Backup and Recovery

Introduction to PostgreSQL



PostgreSQL

AGENDA

- Logical vs physical backup
- Tools: pg_dump, pg_dumpall
- Recovering from a backup
- pg_restore, pg_basebackup, pg_backrest



BACKUPS

- Three basic types of backups
- SQL dump
 - Create a set of SQL files that, when executed, recreate the server data
- File system level backup
 - Makes a physical copy of the file system
- Continuous archiving
 - Combination of a file system back up and WALL file backup
- There various tools that can perform each of these

PHYSICAL BACKUP

- Copying the physical files that PostgreSQL uses to store data
 - Includes the data files, write-ahead logs (WAL), configuration files, and other system files.
- File System Backup:
 - A physical copy of the data directory using tools like rsync, tar, or cp.
 - The database must be in a consistent state
 - Can be done by stopping the server
 - Or by taking the backup while the database is in a backup mode using `pg_start_backup()` and `pg_stop_backup()`
- `pg_basebackup`:
 - A built-in tool that creates a physical backup of the entire database cluster. It can be run while the server is running and supports streaming WAL files for point-in-time recovery.

PHYSICAL BACKUP

- `pg_start_backup()`
 - Used to prepare the database for a file system-level backup by putting it into backup mode.
 - Creates a backup label file and triggers a checkpoint, ensuring that all data files are flushed to disk and the database's state is consistent at the start of the backup.
 - `“SELECT pg_start_backup('backup_label');”`
 - *backup_label*: A text label used to identify the backup session
 - A checkpoint is created to ensure all data files are in a consistent state.
 - A backup label file is created in the data directory, marking the start of the backup.
 - PostgreSQL begins to track which WAL files are necessary to recover the database from the backup point.

PHYSICAL BACKUP

- `pg_stop_backup()`
 - Used to end the backup mode that was started with `pg_start_backup()`.
 - Finalizes the backup and removes the backup label file
 - Records the end of the backup in the WAL stream, ensuring that all changes since the backup began are properly accounted for
 - A backup history file is created in the WAL archive, which logs the backup start and end times and required WAL files.
 - PostgreSQL is returned to normal operation, and the system continues processing WAL as usual.

PHYSICAL BACKUP

- `pg_basebackup`
 - Performs a physical backup by copying the entire data directory, including all database files, configuration files, and the Write-Ahead Log (WAL) files necessary to ensure the consistency of the backup
 - Automatically manages the WAL files needed to restore the backup to a consistent state.
 - If used with WAL archiving, it supports point-in-time recovery (PITR).
 - Can be run while the server is online, allowing backups to be taken without stopping the database.
 - Supports parallel mode, allowing multiple data streams to be used during the backup process to speed up the process for large databases.
 - Supports compression of the backup files to reduce disk space usage and improve transfer speed.

PHYSICAL BACKUP

- “pg_basebackup -D /path/to/backup_directory -F format -h host -p port -U user -W”
 - -D /path/to/backup_directory: directory where the backup will be stored.
 - -F format: p for plain directory, t for tar.
 - -h host: The hostname of the PostgreSQL server.
 - -p port: The port number where PostgreSQL is running.
 - -U user: The username to connect with, typically a replication user.
 - -W: Prompts for a password.
 - *pg_basebackup -D /var/lib/postgresql/backups/ -F t -h localhost -p 5432 -U pgadmin -W*

PHYSICAL BACKUP

- Consistency:
 - Physical backups ensure a consistent state of the database at the file level.
 - When used with WAL archiving, they support point-in-time recovery.
- Performance:
 - Generally faster and more efficient for large databases because they operate at the file level.
- Restoration:
 - Restoration involves copying the backed-up files back into the PostgreSQL data directory and starting the server.
 - Point-in-time recovery can be achieved using archived WAL segments.
- Ideal for disaster recovery, full cluster backup, and high-availability setups (e.g., setting up replicas).

PHYSICAL BACKUP

- Advantages
 - Efficient for large datasets.
 - Supports point-in-time recovery when combined with WAL archiving.
 - Suitable for creating replicas in high-availability configurations.
- Disadvantages
 - Less flexible for selective restoration of individual tables or databases.
 - Larger backup size compared to logical backups.

RESTORING PHYSICAL BACKUP

- Stop the PostgreSQL server
 - Restoring involves replacing the data directory, and running processes can interfere with this.
- Prepare the data directory
 - Ensure that the directory where the backup will be restored is empty.
- Extract or copy the backup
 - If the backup is archived, extract to the data directory
“tar -xzf /path/to/backup/base.tar.gz -C /var/lib/postgresql/16/main/”
 - If the backup is not archived, copy it to the data directory
“sudo cp -R /path/to/backup/ /var/lib/postgresql/16/main/”*
- Set correct permissions
 - Data directory and all its contents should be owned by the PostgreSQL user, usually postgres.

RESTORING PHYSICAL BACKUP

- Configure PostgreSQL for recovery
 - Create a *recovery.signal* file in the data directory.
 - This tells PostgreSQL that it should start in recovery mode.
- Start the PostgreSQL server
 - Monitor the logs to ensure that the recovery process completes successfully.

PHYSICAL BACKUP

- pgBackRest
 - Open-source backup and restore tool for PostgreSQL
 - Designed for full and incremental physical backups, point-in-time recovery, and replication.
 - Allows for efficient backups by only saving changes since the last full backup (incremental) or since the last backup of any type (differential).
 - Supports PITR using archived WAL files, which allows restoring the database to a specific point in time.
 - Supports compressing and encrypting backups
 - Uses multiple processes to speed up backup and restore operations, making it efficient for large databases.

POINT IN TIME RECOVERY

- Point-in-Time Recovery (PITR) allows you to restore a database to a specific moment in time.
 - commonly used for disaster recovery, undoing unintended changes, or recovering from corruption.
 - PITR combines the use of base backups and Write-Ahead Log (WAL) files to perform the recovery
 - The steps are outlined in the notes

LOGICAL BACKUP

- Involve exporting the database's schema and data in a human-readable format (e.g., SQL scripts).
 - Captures the logical content of the database, such as tables, indexes, sequences, and their data, without concern for the underlying file structure.
- `pg_dump`:
 - Utility that exports data and schema definitions of individual tables, databases, or clusters into SQL scripts or other formats like plain text, custom, or directory formats.
- `pg_dumpall`:
 - Utility that exports all databases in a cluster, including global objects like roles and tablespaces.

LOGICAL BACKUP

- `pg_dump`:
 - Performs logical backups by generating SQL scripts or other formats that describe the database objects and their data, which can be executed to recreate the database.
- Multiple output Formats:
 - *Plain Text Format (-Fp)*: Generates a plain SQL script with CREATE, INSERT, and COPY commands.
 - *Easy to read and edit but requires manual parsing for selective restores.*
 - *Custom Format (-Fc)*: A compressed, non-text format that allows selective restoration of database objects using `pg_restore`.
 - *This format is often preferred for flexibility.*
 - *Directory Format (-Fd)*: Dumps the database into a directory with separate files for each table and its data, enabling parallel processing and easy inspection of individual components.
 - *Tar Format (-Ft)*: Outputs the dump in a tar archive. Suitable for easy transport and storage, similar to the directory format but within a single file.

LOGICAL BACKUP

- Selectively back up specific tables, schemas, or data without dumping the entire database
 - Backing up a table
pg_dump -t my_table -Fc -f my_table.dump mydatabase
 - Backing up a schema
pg_dump -n my_schema -Fc -f my_schema.dump mydatabase
 - Can back up tables, views, indexes, sequences, functions, triggers, and other database objects
 - It preserves object ownership, permissions, and other metadata essential for a complete database restoration.

LOGICAL RESTORE

- Restoring with pg_restore
 - When using custom or directory formats, the backups can be restored using pg_restore
 - Provides options for selective restoration of objects and data
 - *-d: Specifies the target database.*
 - *-1: Runs the restore in a single transaction, which is useful for rollback on failure.*
 - *pg_restore -U postgres -d mydatabase -1 mydatabase.dump*
- Selective Restore
 - Restore only specific tables or schemas from a backup:
 - *pg_restore -U postgres -d mydatabase -t my_table mydatabase.dump*

LOGICAL BACKUP

- Database Migrations
 - Export databases for migration between PostgreSQL versions
- Selective Restores
 - Restore individual tables, schema, or specific parts of a database without affecting other objects.
- Development and Testing
 - Create logical backups for setting up development and testing environments quickly.
- Regular Backup Strategy
 - Regular, scheduled logical backups protect against data loss.

LOGICAL BACKUP

- Disadvantages
 - For very large databases, pg_dump can be slower than physical backups, especially if using plain text format.
 - While pg_dump provides consistent snapshots, it does not support point-in-time recovery directly.
 - Running pg_dump can consume significant CPU and I/O resources

LOGICAL BACKUP

- `pg_dumpall`
 - Used to back up an entire database cluster, including all databases, global objects, roles, tablespaces, and configuration settings.
 - Performs a logical backup of everything within a cluster to support full cluster backups and migrations.
- Cluster-Wide Backups
 - `pg_dumpall` creates a comprehensive backup of all databases in a PostgreSQL cluster, including global objects that are not tied to any specific database.
 - Includes global objects such as roles (users), group roles, tablespaces, and database configurations, which are not included when using `pg_dump` alone.
- Output Format:
 - Outputs a plain text SQL script. Unlike `pg_dump`, it does not support custom, directory, or tar formats.
 - Useful reading, editing, and applying during restore.

LOGICAL BACKUP

- Uses pg_dump Internally:
 - pg_dumpall uses pg_dump to back up each database within the cluster.
 - It sequentially dumps each database and combines these dumps into a single output file.
- Basic Usage
 - “pg_dumpall [OPTIONS]”
 - Options
 - *-U: Specifies the user to connect as.*
 - *-h: Specifies the host of the PostgreSQL server.*
 - *-p: Specifies the port number.*
 - *-f: Specifies the output file for the dump.*
 - *-I: Specifies a file containing the list of databases to be dumped.*
 - “pg_dumpall -U postgres -h localhost -p 5432 -f backup.sql”

LOGICAL BACKUP

- Options for pg_dumpall
 - -g or --globals-only:
 - *Dumps only global objects (roles, tablespaces) without the individual databases.*
 - *“pg_dumpall -U postgres --globals-only -f globals_backup.sql”*
 - -r or --roles-only:
 - *Dumps only the roles (users) defined in the cluster.*
 - *“pg_dumpall -U postgres --roles-only -f roles_backup.sql”*
 - --clean:
 - *Includes commands to drop the objects before recreating them, useful for restoring into a cluster where objects might already exist.*
 - *“pg_dumpall -U postgres --clean -f clean_full_cluster_backup.sql”*

LOGICAL BACKUP

- Full Cluster Backups:
 - Use `pg_dumpall` for full backups of the whole cluster, including all databases and global configurations
 - For disaster recovery scenarios where the entire cluster needs to be restored.
- Cluster Migration:
 - Migrate the entire, including roles and tablespaces, from one server to another
 - Or migrate the cluster to different versions of PostgreSQL.
- Consistency Across Databases:
 - Ensure consistency when backing up multiple related databases within the same cluster, since `pg_dumpall` captures a snapshot of the entire cluster.
- Backing Up Global Objects:
 - Use to back up global objects like roles and tablespaces, which are not included in individual `pg_dump` backups.

LOGICAL BACKUP

- Output Size:
 - The plain text output of `pg_dumpall` can be very large – mitigated with compression tools
- Performance:
 - Each database processed sequentially
 - The backup process can take a long time, especially for large clusters.
 - Schedule backups during off-peak hours to minimize impact.
- Restore Process:
 - Restoring executes a large SQL script, which can be slow for large datasets.
- No Custom Formats:
 - Does not support custom or parallelizable formats which may impact performance

LOGICAL RESTORE

- Restoring from a pg_dumpall backup
 - “psql -U postgres -f backup.sql”
 - This executes the SQL in the backup file
 - Recreates the entire cluster, including roles, configurations, and all databases.

LAB 7-1

- The lab description and documentation is in the Lab directory in the class repository



End Module



PostgreSQL